

1. Imports

```
In [14]: import pandas as pd
import pandas_ta as ta
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import mplfinance as mpf
```

```
In [15]: # Read the datafile
df = pd.read_csv('eod_data.csv')
```

```
In [16]: df.head()
```

```
Out[16]:
```

	date	open	high	low	close	volume
0	2007-04-02	5267.70	5267.70	4946.55	4959.65	12535000
1	2007-04-03	4961.30	5048.60	4921.80	4995.15	14984000
2	2007-04-04	4997.05	5108.15	4997.05	5031.75	9251000
3	2007-04-05	5035.90	5151.25	4985.30	5129.20	10120000
4	2007-04-09	5129.05	5335.95	5129.05	5323.80	9708000

```
In [17]: # Convert 'date' column to datetime and set as index
df.set_index(pd.to_datetime(df['date']), inplace=True)
```

```
# Drop the original 'date' column
df.drop(columns=['date'], inplace=True)
```

```
In [18]: df.head()
```

```
Out[18]:
```

	open	high	low	close	volume
date					
2007-04-02	5267.70	5267.70	4946.55	4959.65	12535000
2007-04-03	4961.30	5048.60	4921.80	4995.15	14984000
2007-04-04	4997.05	5108.15	4997.05	5031.75	9251000
2007-04-05	5035.90	5151.25	4985.30	5129.20	10120000
2007-04-09	5129.05	5335.95	5129.05	5323.80	9708000

```
In [19]: df.tail(20)
```

```
Out[19]:
```

	open	high	low	close	volume
date					
2024-02-09	44986.7500	45718.1500	44859.1500	45634.5500	275980000
2024-02-12	45664.3000	45748.5000	44633.8500	44882.2500	268820000
2024-02-13	45056.8000	45750.4000	44819.5500	45502.4000	245860000
2024-02-14	45014.6500	46170.4500	44860.7500	45908.3000	279950000
2024-02-15	46027.1000	46297.7000	45590.2000	46218.9000	277740000
2024-02-16	46454.3000	46693.4000	46264.4000	46384.8500	291380000
2024-02-19	46554.9000	46717.4000	46317.7000	46535.5000	158390000
2024-02-20	46444.9000	47136.7500	46367.8000	47094.2000	171890000
2024-02-21	47363.4000	47363.4000	46886.9500	47019.7000	198950000
2024-02-22	46934.5500	47024.0500	46426.8500	46919.8000	177180000
2024-02-23	47060.7000	47245.3500	46723.1500	46811.7500	200520000
2024-02-26	46615.8500	46893.1500	46513.5500	46576.5000	223680000
2024-02-27	46480.2000	46722.2500	46324.9000	46588.0500	157260000
2024-02-28	46640.9000	46754.5500	45852.5500	45963.1500	167270000
2024-02-29	45881.4500	46329.6500	45661.7500	46120.9000	544730000
2024-03-01	46218.0000	47342.2500	46218.0000	47286.8984	172500
2024-03-02	47342.1000	47392.9500	47240.3500	47297.5000	12240
2024-03-04	47318.5000	47529.6016	47191.6484	47456.1016	158100
2024-03-05	47265.6992	47737.8516	47196.7500	47581.0000	176900
2024-03-06	47451.6484	48161.2500	47442.2500	47965.3984	301700

```
In [20]: df.shape
```

```
Out[20]: (4188, 5)
```

```
In [20]:
```

2. Add Technical Indicators

```
In [21]: # Compute Simple Moving Averages
df['sma_10'] = df['close'].rolling(window=10).mean()
df['sma_20'] = df['close'].rolling(window=20).mean()
df['sma_50'] = df['close'].rolling(window=50).mean()

# Compute Exponential Moving Averages
df['ema_50'] = df.ta.ema(length=50, close='close')
df['ema_200d'] = df.ta.ema(length=200, close='close')

# Compute high and low of the last 250 days
df['hi_250d'] = df['high'].rolling(window=250).max()
df['lo_250d'] = df['low'].rolling(window=250).min()

# Compute previous close
df['prev_close'] = df['close'].shift(1)

# Compute change percentage
df['change_p'] = df['close'].pct_change() * 100

# Compute average volume
df['avgvol_14d'] = df['volume'].rolling(window=14).mean()
df['avgvol_50d'] = df['volume'].rolling(window=50).mean()
df['avgvol_200d'] = df['volume'].rolling(window=200).mean()

# Compute RSI
df['rsi'] = ta.rsi(df['close'], length=14)

# Compute MACD
macd_indicators = df.ta.macd(fast=12, slow=26, signal=9)
df = pd.concat([df, macd_indicators], axis=1)
```

```
df['vwap'] = df.ta.vwap(high='high', low='low', close='close', volume='volume')

# Add Stochastic Oscillator
df[['stoch_k', 'stoch_d']] = df.ta.stoch(high='high', low='low', close='close')

# Add ADX
adx = df.ta.adx(high='high', low='low', close='close')
df = pd.concat([df, adx], axis=1)

# Add Accumulation/Distribution Line
df['ad_line'] = df.ta.ad(high='high', low='low', close='close', volume='volume')
```

3. Data Exploration

```
In [22]: df.head(1000)
```

Out[22]:

	open	high	low	close	volume	sma_10	sma_20	sma_50	ema_50	ema_200d	...	MACD_12
date												
2007-04-02	5267.70	5267.70	4946.55	4959.65	12535000	NaN	NaN	NaN	NaN	NaN	...	
2007-04-03	4961.30	5048.60	4921.80	4995.15	14984000	NaN	NaN	NaN	NaN	NaN	...	
2007-04-04	4997.05	5108.15	4997.05	5031.75	9251000	NaN	NaN	NaN	NaN	NaN	...	
2007-04-05	5035.90	5151.25	4985.30	5129.20	10120000	NaN	NaN	NaN	NaN	NaN	...	
2007-04-09	5129.05	5335.95	5129.05	5323.80	9708000	NaN	NaN	NaN	NaN	NaN	...	
...	
2011-04-11	11591.00	11637.00	11551.75	11597.80	12854000	11700.280	11334.4050	10938.301	11185.555455	10940.248520	...	236.7
2011-04-13	11531.45	11856.35	11490.55	11839.50	15317000	11725.980	11390.0900	10962.254	11211.200339	10949.196295	...	240.6
2011-04-15	11810.10	11950.70	11707.60	11735.65	17418000	11723.025	11428.9925	10987.343	11231.766993	10957.021705	...	232.6
2011-04-18	11754.20	11887.10	11525.05	11539.75	18486000	11706.455	11463.3825	11008.735	11243.844758	10962.819997	...	208.1
2011-04-19	11562.95	11631.15	11482.15	11586.75	20322000	11706.035	11506.9250	11027.058	11257.292022	10969.028256	...	190.2

1000 rows × 28 columns

```
In [23]: df.tail()
```

Out[23]:

	open	high	low	close	volume	sma_10	sma_20	sma_50	ema_50	ema_200d
date										
2024-03-01	46218.0000	47342.2500	46218.0000	47286.8984	172500	46691.64484	46189.72742	46686.599968	46279.585977	44903.385902
2024-03-02	47342.1000	47392.9500	47240.3500	47297.5000	12240	46767.84484	46263.32492	46683.643968	46319.504174	44927.207933
2024-03-04	47318.5000	47529.6016	47191.6484	47456.1016	158100	46804.03500	46351.59000	46675.963000	46364.076622	44952.371054
2024-03-05	47265.6992	47737.8516	47196.7500	47581.0000	176900	46860.16500	46439.71500	46677.746000	46411.799108	44978.526566
2024-03-06	47451.6484	48161.2500	47442.2500	47965.3984	301700	46964.72484	46587.38492	46682.556968	46472.724570	45008.246683

5 rows × 28 columns

```
In [24]: print(f"Date is from dates between {df.index[0]} and {df.index[-1]}")
```

Date is from dates between 2007-04-02 00:00:00 and 2024-03-06 00:00:00

```
In [25]: # Check the number of duplicates
df.isna().sum()
```

```
Out[25]: open          0
         high          0
         low           0
         close         0
         volume        0
         sma_10         9
         sma_20        19
         sma_50        49
         ema_50         49
         ema_200d      199
         hi_250d       249
         lo_250d       249
         prev_close     1
         change_p       1
         avgvol_14d     13
         avgvol_50d     49
         avgvol_200d    199
         rsi            14
         MACD_12_26_9   25
         MACDh_12_26_9  33
         MACDs_12_26_9  33
         vwap           0
         stoch_k        15
         stoch_d        17
         ADX_14         27
         DMP_14         14
         DMN_14         14
         ad_line        0
         dtype: int64
```

```
In [26]: df.columns
```

```
Out[26]: Index(['open', 'high', 'low', 'close', 'volume', 'sma_10', 'sma_20', 'sma_50',
               'ema_50', 'ema_200d', 'hi_250d', 'lo_250d', 'prev_close', 'change_p',
               'avgvol_14d', 'avgvol_50d', 'avgvol_200d', 'rsi', 'MACD_12_26_9',
               'MACDh_12_26_9', 'MACDs_12_26_9', 'vwap', 'stoch_k', 'stoch_d',
               'ADX_14', 'DMP_14', 'DMN_14', 'ad_line'],
              dtype='object')
```

```
In [27]: # Create the response columns, using the predictor values from the next day
```

```
df['op_nxt'] = df['open'].shift(-1)
df['hi_nxt'] = df['high'].shift(-1)
df['lo_nxt'] = df['low'].shift(-1)
df['cl_nxt'] = df['close'].shift(-1)
```

```
# Move these columns to be beginning of the dataframe
```

```
columns = ['op_nxt', 'hi_nxt', 'lo_nxt', 'cl_nxt', 'open', 'high', 'low', 'close', 'volume',
           'sma_10', 'sma_20', 'sma_50', 'ema_50', 'ema_200d', 'hi_250d', 'lo_250d',
           'prev_close', 'change_p', 'avgvol_14d', 'avgvol_50d', 'avgvol_200d', 'rsi',
           'MACD_12_26_9', 'MACDh_12_26_9', 'MACDs_12_26_9', 'vwap',
           'stoch_k', 'stoch_d', 'ADX_14', 'DMP_14', 'DMN_14', 'ad_line']
```

```
df = df[columns]
```

```
In [28]: df.head(10)
```

```
Out[28]:
```

	op_nxt	hi_nxt	lo_nxt	cl_nxt	open	high	low	close	volume	sma_10	...	MACD_12_26_9	MACDh_12_26_9
date													
2007-04-02	4961.30	5048.60	4921.80	4995.15	5267.70	5267.70	4946.55	4959.65	12535000	NaN	...	NaN	NaN
2007-04-03	4997.05	5108.15	4997.05	5031.75	4961.30	5048.60	4921.80	4995.15	14984000	NaN	...	NaN	NaN
2007-04-04	5035.90	5151.25	4985.30	5129.20	4997.05	5108.15	4997.05	5031.75	9251000	NaN	...	NaN	NaN
2007-04-05	5129.05	5335.95	5129.05	5323.80	5035.90	5151.25	4985.30	5129.20	10120000	NaN	...	NaN	NaN
2007-04-09	5335.40	5349.75	5229.75	5325.50	5129.05	5335.95	5129.05	5323.80	9708000	NaN	...	NaN	NaN
2007-04-10	5325.30	5374.25	5261.30	5276.65	5335.40	5349.75	5229.75	5325.50	9197000	NaN	...	NaN	NaN
2007-04-11	5241.15	5241.15	5176.45	5215.15	5325.30	5374.25	5261.30	5276.65	5527000	NaN	...	NaN	NaN
2007-04-12	5234.50	5404.10	5230.90	5362.10	5241.15	5241.15	5176.45	5215.15	5884000	NaN	...	NaN	NaN
2007-04-13	5400.30	5465.50	5400.30	5435.10	5234.50	5404.10	5230.90	5362.10	10973000	NaN	...	NaN	NaN
2007-04-16	5457.20	5522.15	5387.50	5455.45	5400.30	5465.50	5400.30	5435.10	7437000	5205.405	...	NaN	NaN

10 rows × 32 columns

```
In [29]: # Check the datatypes
print(df.dtypes)
```

```
op_nxt      float64
hi_nxt      float64
lo_nxt      float64
cl_nxt      float64
open        float64
high        float64
low         float64
close       float64
volume      int64
sma_10      float64
sma_20      float64
sma_50      float64
ema_50      float64
ema_200d    float64
hi_250d     float64
lo_250d     float64
prev_close  float64
change_p    float64
avgvol_14d  float64
avgvol_50d  float64
avgvol_200d float64
rsi         float64
MACD_12_26_9 float64
MACDh_12_26_9 float64
MACDs_12_26_9 float64
vwap        float64
stoch_k     float64
stoch_d     float64
ADX_14      float64
DMP_14      float64
DMN_14      float64
ad_line     float64
dtype: object
```

```
In [30]: # Check the missing values
df.isna().sum()
```

```

Out[30]:op_nxt          1
        hi_nxt          1
        lo_nxt          1
        cl_nxt          1
        open            0
        high            0
        low             0
        close           0
        volume          0
        sma_10          9
        sma_20         19
        sma_50         49
        ema_50         49
        ema_200d       199
        hi_250d       249
        lo_250d       249
        prev_close      1
        change_p        1
        avgvol_14d      13
        avgvol_50d      49
        avgvol_200d    199
        rsi             14
        MACD_12_26_9    25
        MACDh_12_26_9   33
        MACDs_12_26_9   33
        vwap            0
        stoch_k         15
        stoch_d         17
        ADX_14          27
        DMP_14          14
        DMN_14          14
        ad_line         0
        dtype: int64

In [31]: df.shape

Out[31]: (4188, 32)

In [32]: # Drop all the rows with missing values
        df = df.dropna()

In [33]: df.shape

Out[33]: (3938, 32)

In [34]: df.isna().sum()

Out[34]:op_nxt          0
        hi_nxt          0
        lo_nxt          0
        cl_nxt          0
        open            0
        high            0
        low             0
        close           0
        volume          0
        sma_10          0
        sma_20          0
        sma_50          0
        ema_50          0
        ema_200d        0
        hi_250d         0
        lo_250d         0
        prev_close      0
        change_p        0
        avgvol_14d      0
        avgvol_50d      0
        avgvol_200d     0
        rsi             0
        MACD_12_26_9    0
        MACDh_12_26_9   0
        MACDs_12_26_9   0
        vwap            0
        stoch_k         0
        stoch_d         0
        ADX_14          0
        DMP_14          0
        DMN_14          0
        ad_line         0
        dtype: int64

```

```

In [35]: # Save the dataframe as csv
df.to_csv('data_ml.csv', index=True)

In [36]: # Define the response and the predictors
response = df['cl_nxt'] # following day's close
predictors = df[['open', 'high', 'low', 'volume', 'close',
                 'sma_10', 'sma_20', 'sma_50', 'ema_50', 'ema_200d', 'hi_250d', 'lo_250d',
                 'prev_close', 'change_p', 'avgvol_14d', 'avgvol_50d', 'avgvol_200d', 'rsi',
                 'MACD_12_26_9', 'MACDh_12_26_9', 'MACDs_12_26_9', 'vwap',
                 'stoch_k', 'stoch_d', 'ADX_14', 'DMP_14', 'DMN_14', 'ad_line']] # ohlc, volume, & technci

In [37]: # First, concatenate predictors and response in the same DataFrame
df_combined = pd.concat([predictors, response], axis=1)
df_combined.shape

Out[37]: (3938, 29)

In [:]: # Plot the candlestick
figsize = (20,8)
df_candlestick = df[['open', 'high', 'low', 'volume', 'close']]
mpf.plot(df_candlestick, type='candle', style='yahoo', volume=True, figsize=figsize)

```

Candlesticks on a chart represent price movements within a set period, displaying the open, high, low, and close values, with the body's color indicating whether the closing price was higher (usually green) or lower (usually red) than the opening price. Trends are discerned from the direction and patterns of these candlesticks over time, indicating an upward, downward, or sideways market movement. Volume bars complement candlesticks by showing the quantity of an asset traded during the corresponding period; tall bars indicate high trading activity, which can validate the strength of a price move, while short bars suggest less trading activity and potentially less conviction in the price trend.

This Bank Nifty chart presents an upward trajectory from 2008 to 2022 with discernible periods of highs and lows. The lows or dips, where there is a noticeable decline in the index value, can be spotted at specific intervals which could be associated with broader economic downturns or sector-specific challenges. One such significant dip appears to occur around 2020, which aligns with the global financial impact of the COVID-19 pandemic—a period known for its high market volatility and uncertainty.

Regarding trading volume, we see peaks that often correspond with the index's price fluctuations. For instance, increased volume during the lows suggests heightened trading activity, which often occurs when investors react to market stress by selling off assets, while elevated volumes during the highs may reflect increased buying activity as investor confidence grows and they re-enter the market to capitalize on the anticipated recovery and growth. These volume peaks provide a narrative of investor sentiment, with high volumes in downturns indicating potential capitulation or high selling pressure, and high volumes in upswings suggesting strong buying interest.

```

In [:]: # Plot the heatmap

# Compute the correlation matrix
corr = df_combined.corr()

# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(20, 20))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(230, 20, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, center=0, annot=True,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})

plt.show()

```

In the heatmap, the indicators 'open', 'high', 'low', 'volume', and the moving averages 'sma_10', 'sma_20', 'ema_50', and 'ema_200d' have high positive correlation coefficients, mostly close to +1, represented by the deep red color. This indicates that these variables typically move together; when one goes up, the others tend to go up as well, and vice versa. This is expected as they are all directly related to the price action of a security.

On the other hand, indicators like 'DMN_14', 'ad_line' (Advance/Decline Line), and 'd_nxt' appear to have less consistent relationships with the other variables. For instance, 'DMN_14' shows strong negative correlations (blue squares) with several of the price-related indicators, suggesting that when the price indicators are increasing, 'DMN_14' tends to decrease, and this can be characteristic of the indicator showing strength in downward price movements. The 'ad_line' shows a mix of positive and negative correlations with other indicators but tends to be less strongly correlated overall, indicating that its movements are not as closely tied to price changes.

The 'cl_nxt' variable shows a range of correlations with different financial indicators. For most indicators, such as 'open', 'high', 'low', 'volume', and various moving averages like 'sma_10', 'sma_20', 'ema_50', and 'ema_200d', the correlation coefficients are near 0, denoting a very weak or no linear relationship. This implies that these indicators from the current or previous days do not consistently predict the next day's closing price. However, there are a few indicators with a stronger relationship; for instance, 'cl_nxt' shows a moderately negative correlation with 'DMN_14', as indicated by a lighter blue square. This suggests that the previous day's 'DMN_14' values have some degree of inverse association with the next day's closing price, though it's not strong enough to be highly predictive.

```
In[: # Identifying Strong Correlations
correlation_matrix = predictors.corrwith(response)

# Identify predictors that have an absolute correlation greater than 0.5 with response
significant_correlations = (correlation_matrix[correlation_matrix.abs() > 0.5]).sort_values(ascending=

print(significant_correlations)
```

The correlations listed with 'cl_nxt' are extremely high, especially with 'close', 'vwap' (volume-weighted average price), 'high', 'low', and 'open', all above 0.999. This indicates an almost perfect linear relationship; as these variables change, the next day's closing price is likely to move in the same direction nearly one-to-one.

Moving averages like 'sma_10', 'sma_20', 'ema_50', 'sma_50', and 'ema_200d' also show very high correlations, decreasing slightly as the number of days in the moving average increases, which suggests that while these are still highly predictive of the next day's close, the relationship is slightly less direct due to the smoothing effect of these indicators over more extended periods.

'hi_250d' and 'lo_250d' represent the 250-day highs and lows, and they also correlate highly, but less so than the daily indicators, reflecting that historical extremes have a lesser, yet still strong, influence on the next day's closing price.

'ad_line' or advance-decline line, a cumulative measure of the number of advancing and declining issues on an exchange, shows a strong correlation but less so than price-related indicators, hinting that broader market movements have a significant, but less immediate, impact on the next day's closing price.

Average volumes over longer periods ('avgvol_200d', 'avgvol_50d', and 'avgvol_14d') show moderate correlations, indicating that higher trading volumes can influence the next day's price, possibly through sustained buying or selling pressure, but with less predictive power than price movements.

Lastly, 'volume' has the lowest correlation of the listed indicators, suggesting that while there is some relationship between trading volume on a given day and the next day's closing price, the connection is weaker, potentially due to daily volume being influenced by short-term events that may not have a lasting impact on price.

```
In[: import matplotlib.pyplot as plt
df_combined.hist(figsize=(20,20))
plt.show()
```

The 'open', 'high', 'low', 'close', 'sma_10', 'sma_20', 'sma_50', 'ema_50', 'ema_200d', 'hi_250d', 'lo_250d', 'prev_close', and 'cl_nxt' histograms have a right-skewed distribution, indicating a higher frequency of lower values and fewer high values.

The 'volume', 'avgvol_50d', 'avgvol_200d', and 'avgvol_14d' histograms display a highly right-skewed distribution, suggesting a concentration of data points towards the lower end of the volume scale with very few instances of extremely high volume.

The 'change_p' histogram seems to be normally distributed around 0, with tails extending to both positive and negative changes, indicating that price changes fluctuate symmetrically around no change.

The 'rsi', 'stock_k', 'stock_d', 'DMP_14', and 'DMN_14' histograms appear to have a more uniform or slightly bimodal distribution, indicating that the data points are spread out across the range of values, with some concentrations in specific intervals.

The 'MACD_12_26_9' and 'MACDh_12_26_9' histograms show a distribution centered around zero, with the MACD histogram being slightly left-skewed, suggesting more frequent occurrence of negative values.

The 'vwap', 'ad_line', and 'ADX_14' histograms are moderately skewed, indicating that while there's a range of values, there's a tendency toward one end of the spectrum.

```
In[: # Plot the boxplots
features = df_combined.columns
fig, axs = plt.subplots(nrows=6, ncols=5, figsize=(20, 24)) # adjust the size as needed
axs = axs.flatten() # to iterate over the grid easily

for i, column in enumerate(features):
    sns.boxplot(x=df_combined[column], ax=axs[i])
    axs[i].title.set_text(column)

# In case of less than 30 features, delete unused subplots
if len(features) < 30:
    for i in range(len(features), 30):
        fig.delaxes(axs[i])

plt.tight_layout()
plt.show()
```


For 'open', 'high', 'low', 'close', 'sma_10', 'sma_20', 'sma_50', 'ema_50', 'ema_200d', 'hi_250d', 'lo_250d', 'prev_close', 'wvap', 'stock_k', and 'stock_d', the box plots show a relatively symmetrical distribution with the median line near the center of the box, suggesting a more or less even distribution of data around the median.

The 'volume', 'avgvol_50d', 'avgvol_200d', and 'avgvol_14d' indicators have box plots with a line (median) closer to the bottom of the box, indicating a right-skewed distribution, with a few outliers indicating instances of extremely high volume.

'change_p', 'MACD_12_26_9', 'MACDh_12_26_9', and 'ADX_14' display box plots with medians close to zero but with various spreads and outliers, indicating occasional extreme values or fluctuations from the typical range.

The 'rsi' indicator's box plot shows the median closer to the upper quartile, which might indicate a distribution that is slightly skewed towards higher values.

Outliers are shown as individual points outside the 'whiskers' of the box plots, which represent 1.5 times the interquartile range (the distance between Q1 and Q3). These outliers suggest that there are values that deviate significantly from the rest of the distribution.

Summary Data Exploration:

The data exploration reveals that Bank Nifty's price-related indicators ('open', 'high', 'low', 'close', and various moving averages) generally move in tandem, as evidenced by the near-perfect correlation coefficients with 'cl_nxt', the next day's closing price. These indicators also show right-skewed distributions, suggesting a higher occurrence of lower values. 'Volume' and average volume indicators exhibit right-skewed distributions in both box plots and histograms, indicating that most trading days have low to moderate volume, with occasional spikes reflecting periods of high trading activity. Such spikes in volume are often associated with significant price movements, hinting at reactive trading during times of market stress or heightened investor interest.

The histograms for 'change_p' and the MACD indicators center around zero, revealing that day-to-day price changes and momentum oscillations are normally distributed, which indicates a balanced dynamic of price fluctuations around a mean value. The RSI's distribution leans towards higher values, suggesting a period of generally positive momentum for Bank Nifty.

The overall trend of Bank Nifty from 2008 to 2022 is positive, with noticeable dips correlating to broader market downturns, such as the 2020 COVID-19 market impact. Recovery post-dips suggests resilience in the banking index, and the peaks in trading volume around these lows and subsequent highs signal active market participation, with investors likely selling during downturns and buying during upturns.

```
In[: !pip install nbconvert
```

```
In[: !jupyter nbconvert --to html data_exploration.ipynb
```

```
In[: !wkhtmltopdf ./data_exploration.html data_exploration.pdf
```